

Lecture 3.1

Supervised learning: Bayesian and linear classifiers



Universitat
de les Illes Balears

Departament
de Ciències Matemàtiques
i Informàtica

11752 Aprendizaje Automático
11752 Machine Learning
Máster Universitario
en Sistemas Inteligentes

Alberto ORTIZ RODRÍGUEZ

- Introduction
- Bayesian classification
- Estimation of probability density functions
- Linear discriminant functions and the perceptron algorithm

- **Supervised classification**

- It is about classifying a new sample in the correct class, having initially designed a classifier from the data available in a training set, in which, in particular, the **samples are labeled** with the class to which they belong.

- Introduction
- Bayesian classification
- Estimation of probability density functions
- Linear discriminant functions and the perceptron algorithm

Bayesian classification

- The **goal** is to classify a new sample in the **most likely class**
 - Given a **classification task** in M classes, $\omega_1, \omega_2, \dots, \omega_M$, and a **new sample** x , we deal with:

$p(\omega_i|x), i = 1, 2, \dots, M$ (**probabilities a posteriori**)

- The classifier decides the most likely class based on the **maximum** of the probabilities *a posteriori*:
 - **Bayesian classification rule**

if $p(\omega_i|x) > p(\omega_j|x), \forall j \neq i$, then x is labelled as class i

Bayesian classification


- Review of probability theory:

- **Probability function:** $p : \Omega \rightarrow [0, 1]$

$$v \rightarrow p(v)$$

- p assigns a value to each possible event v on the basis of how often that event occurs
 - Ω can be stated as the set of events corresponding to a certain discrete **random variable** X taking certain values, so that

$$p(A_i) = p(X = x_i)$$



- Of particular relevance:

$$p(\Omega) = \sum_{i=1}^n p(A_i) = 1$$

Bayesian classification

- Review of probability theory:

- **Law of total probability**

Given M events $A_i, i = 1, \dots, M$, such that $\sum_{i=1}^M p(A_i) = 1$,
for any random event:

$$p(B) = \sum_{i=1}^M p(B|A_i)p(A_i)$$

where the probability of B conditioned to the occurrence of event A_i is defined as:

$$p(B|A_i) = \frac{p(B \cap A_i)}{p(A_i)}$$

- **Bayes Rule**

From the definition of conditional probability, given two events A and B :

$$p(B|A)p(A) = p(A|B)p(B)$$

✿ *All this is verified under exactly the same conditions by substituting probabilities by probability density functions (**pdf**'s)*

Bayesian classification


- Bayesian classification: **two-class case** (ω_1, ω_2)
 - Given the **probabilities *a priori*** of both classes $p(\omega_1)$ and $p(\omega_2)$
 - If you do not know $p(\omega_1)$ and $p(\omega_2)$, you can estimate them if necessary:

$$p(\omega_1) \approx \frac{n_1}{n_1+n_2}, \quad p(\omega_2) \approx \frac{n_2}{n_1+n_2}$$

and the **pdf (probability density functions)** of each class

$$p(x|\omega_i), i = 1, 2$$

- If they are unknown, they have to be estimated from the available training data (we will deal with this topic later)



using **Bayes' rule**, it follows that:

$$p(\omega_i|x) = \frac{p(x|\omega_i)p(\omega_i)}{p(x)} = \frac{p(x|\omega_i)p(\omega_i)}{\sum_{j=1}^2 p(x|\omega_j)p(\omega_j)}$$

↑
total probability

Bayesian classification

- Bayesian classification: **two-class case** (ω_1, ω_2)

$$\left. \begin{aligned} p(\omega_i|x) &= \frac{p(x|\omega_i)p(\omega_i)}{p(x)} \\ p(\omega_i|x) &> p(\omega_j|x), \forall j \neq i \end{aligned} \right\} \Rightarrow p(x|\omega_i)p(\omega_i) > p(x|\omega_j)p(\omega_j), \forall j \neq i$$


- If the probabilities *a priori* are equal ($p(\omega_i) = 1/M = 0.5$), then the classification rule becomes dependent only on the **pdfs** of the classes:

$$p(x|\omega_i) > p(x|\omega_j), \forall j \neq i$$

- In the one-dimensional case (1 feature), x_0 is a threshold that partitions the space into two regions, R_1 and R_2


- It is obvious that **classification errors** are unavoidable:

- Sample x can be inside region R_2 and belong to the class ω_1 (same for R_1 and ω_2)



Bayesian classification

- Bayesian classification: **two-class case** (ω_1, ω_2)
 - In the one-dimensional case, the **probability of making a classification error** is given by:



$$\begin{aligned}P_e &= p(x \in R_1 \cap x \text{ is from } \omega_2) + p(x \in R_2 \cap x \text{ is from } \omega_1) \\&= p(x \in R_1 | \omega_2)p(\omega_2) + p(x \in R_2 | \omega_1)p(\omega_1) \\&= p(\omega_2) \int_{R_1} p(x | \omega_2) dx + p(\omega_1) \int_{R_2} p(x | \omega_1) dx \\&= \frac{1}{2} \left(\int_{-\infty}^{x_0} p(x | \omega_2) dx + \int_{x_0}^{+\infty} p(x | \omega_1) dx \right) \\&\quad (\text{if } p(\omega_1) = p(\omega_2) = 0.5)\end{aligned}$$

THEOREM

The Bayesian classifier minimizes the likelihood of classification error

That is to say, if we move x_0 left or right we will increase P_e

Bayesian classification

- Bayesian classification: **two-class case** (ω_1, ω_2)

Proof (optimality of the Bayesian classifier)

On the one hand:

$$\begin{aligned} P_e &= p(x \in R_2 \cap \omega_1) + p(x \in R_1 \cap \omega_2) = p(x \in R_2 | \omega_1)p(\omega_1) + p(x \in R_1 | \omega_2)p(\omega_2) \\ &= p(\omega_1) \int_{R_2} p(x | \omega_1) dx + p(\omega_2) \int_{R_1} p(x | \omega_2) dx \\ &= \boxed{\int_{R_2} p(\omega_1 | x)p(x) dx} + \int_{R_1} p(\omega_2 | x)p(x) dx \end{aligned}$$

On the other hand:

$$\begin{aligned} \int_{\Omega} p(x | \omega_1) dx &= 1 \Rightarrow \int_{\Omega} \frac{p(\omega_1 | x)p(x)}{p(\omega_1)} dx = 1 \\ &\Rightarrow \int_{R_1} p(\omega_1 | x)p(x) dx + \boxed{\int_{R_2} p(\omega_1 | x)p(x) dx} = p(\omega_1) \end{aligned}$$

Therefore:

$$P_e = p(\omega_1) - \int_{R_1} (p(\omega_1 | x) - p(\omega_2 | x))p(x) dx$$

$\Rightarrow P_e$ is minimum if R_1 is defined such that, inside R_1 , $p(\omega_1 | x) > p(\omega_2 | x)$

Bayesian classification

- Bayesian classification:


two-class case (ω_1, ω_2) and **2 features** x_1 and x_2

– Minimize the probability of error is equivalent to **partition the space of features into M regions** (as many as classes)

– In general terms, regions R_i and R_j are separated by a **decision curve** that is described by the equation:

$$p(\omega_i|x) - p(\omega_j|x) = 0$$

- Corresponds to the points of the space in which the probabilities *a posteriori* coincide



Bayesian classification

- Bayesian Classification

- **Example** Let us consider a problem of 2 equiprobable classes ($p(\omega_1) = p(\omega_2) = 0.5$) such that the class pdfs are Gaussians of **variance 0.5** and **means 0 and 1** respectively:

$$p(x|\omega_1) = \frac{1}{\sqrt{\pi}} e^{-x^2}, \quad p(x|\omega_2) = \frac{1}{\sqrt{\pi}} e^{-(x-1)^2} \quad \left(N(\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}} \right)$$

Calculate the optimal threshold x_0 for minimum error probability.

$$x_0 : p(\omega_1|x_0) = p(\omega_2|x_0)$$

$$x_0 : p(x_0|\omega_1)p(\omega_1) = p(x_0|\omega_2)p(\omega_2)$$

$$x_0 : e^{-x_0^2} = e^{-(x_0-1)^2} \Rightarrow x_0^2 = (x_0 - 1)^2 = x_0^2 - 2x_0 + 1 \Rightarrow x_0 = 0.5$$

Bayesian classification

- Bayesian classification for normal distributions

- In the one-dimensional case:

$$p(x|\omega) = \frac{1}{\sqrt{2\pi}\sigma_\omega} e^{-\frac{(x-\mu_\omega)^2}{2\sigma_\omega^2}}$$

- We assume that the pdf of the classes obey the Gaussian L-dimensional distribution:

$$p(x|\omega) = \frac{1}{\sqrt{(2\pi)^L |\Sigma_\omega|}} e^{-\frac{1}{2}(x-\mu_\omega)^T \Sigma_\omega^{-1} (x-\mu_\omega)}$$

- For class ω :

$$\mu_\omega = (\mu_1, \dots, \mu_s, \dots, \mu_L) \quad , \quad \mu_s = \frac{1}{N} \sum_{i=1}^N x_{si}$$


$$\begin{aligned} \Sigma_\omega &= \begin{bmatrix} \sigma_1^2 & \sigma_{12} & \dots & \sigma_{1L} \\ \sigma_{12} & \sigma_2^2 & \dots & \sigma_{2L} \\ \vdots & \vdots & & \vdots \\ \sigma_{L1} & \sigma_{L2} & \dots & \sigma_L^2 \end{bmatrix} \quad , \quad \sigma_{st} = \frac{1}{N-1} \sum_{i=1}^N (x_{si} - \mu_s)(x_{ti} - \mu_t) \\ &= \frac{1}{N-1} \sum_{i=1}^N (x - \mu_\omega)(x - \mu_\omega)^T \end{aligned}$$

- This distribution models properly many cases and is treatable mathematically and computationally, hence its popularity

Bayesian classification

- Bayesian classification for normal distributions
 - Example

$$\mu = (0, 0) \quad \Sigma = \begin{pmatrix} 1 & 1.2 \\ 1.2 & 2 \end{pmatrix}$$



Bayesian classification

- Bayesian classifier for normal distributions

- The goal is to derive the Bayesian classifier for the case

$$p(x|\omega_i) = \frac{1}{\sqrt{(2\pi)^L |\Sigma_i|}} e^{-\frac{1}{2}(x-\mu_i)^T \Sigma_i^{-1} (x-\mu_i)}, i = 1, \dots, M$$

- Due to the exponential form of the pdf, it is preferable to work with the following **discrimination functions** $g_i(x)$, which involve the **monotonous function** $\ln(\cdot)$:

$$g_i(x) = \ln(p(x|\omega_i)p(\omega_i)) = \ln p(x|\omega_i) + \ln p(\omega_i)$$

$$g_i(x) = c_i - \frac{1}{2}(x - \mu_i)^T \Sigma_i^{-1} (x - \mu_i) + \ln p(\omega_i)$$

$$c_i = -\frac{L}{2} \ln 2\pi - \frac{1}{2} \ln |\Sigma_i|$$

- Finally:

$$\begin{aligned} g_i(x) = & -\frac{1}{2} x^T \Sigma_i^{-1} x + \frac{1}{2} x^T \Sigma_i^{-1} \mu_i + \frac{1}{2} \mu_i^T \Sigma_i^{-1} x - \frac{1}{2} \mu_i^T \Sigma_i^{-1} \mu_i \\ & + \ln p(\omega_i) + c_i \end{aligned}$$

Bayesian classification

- Bayesian classifier for normal distributions

- Case of 2 uncorrelated characteristics

$$L = 2, \quad x = (x_1, x_2)^T, \quad \Sigma_i = \begin{pmatrix} \sigma_{i1}^2 & 0 \\ 0 & \sigma_{i2}^2 \end{pmatrix}$$

$$g_i(x) = -\frac{1}{2}x^T \Sigma_i^{-1} x + \frac{1}{2}x^T \Sigma_i^{-1} \mu_i + \frac{1}{2}\mu_i^T \Sigma_i^{-1} x - \frac{1}{2}\mu_i^T \Sigma_i^{-1} \mu_i$$

$$+ \ln p(\omega_i) + c_i$$

$$\Rightarrow g_i(x) = -\frac{1}{2} \left(\frac{x_1^2}{\sigma_{i1}^2} + \frac{x_2^2}{\sigma_{i2}^2} \right) + \left(\frac{\mu_{i1}x_1}{\sigma_{i1}^2} + \frac{\mu_{i2}x_2}{\sigma_{i2}^2} \right) - \frac{1}{2} \left(\frac{\mu_{i1}^2}{\sigma_{i1}^2} + \frac{\mu_{i2}^2}{\sigma_{i2}^2} \right)$$

$$+ \ln p(\omega_i) + c_i$$

$$= \alpha_i x_1^2 + \beta_i x_2^2 + \gamma_i x_1 + \delta_i x_2 + \epsilon_i$$

- The decision rule is now given by the equation $g_i(x) - g_j(x) = 0$

- $L = 2$: ellipse, parabola, hyperbole, etc. – **conic**, rule = 2D curve

- $L = 3$: ellipsoid, paraboloid, hyperboloid, etc. – **quadric**, rule = 3D surface

- $L > 3$: **hiperquadric**

$$\rightarrow g_i(x) > g_j(x) \Rightarrow p(\omega_i|x) > p(\omega_j|x) \Rightarrow x \rightarrow \omega_i$$

QUADRATIC
CLASSIFIER



$$\Rightarrow g_i(x) - g_j(x) = (\alpha_i - \alpha_j)x_1^2 + (\beta_i - \beta_j)x_2^2 + (\gamma_i - \gamma_j)x_1 + (\delta_i - \delta_j)x_2 + (\epsilon_i - \epsilon_j) > 0 \Rightarrow x \rightarrow \omega_i$$

Bayesian classification


- Bayesian classifier for normal distributions
 - Example (equiprobable classes)

$$\mu_1 = (0, 0)^T, \mu_2 = (1, 0)^T, \Sigma_1 = \begin{pmatrix} 0.10 & 0.00 \\ 0.00 & 0.15 \end{pmatrix}, \boxed{\Sigma_2 = \begin{pmatrix} 0.20 & 0.00 \\ 0.00 & 0.25 \end{pmatrix}}$$

$$g_1(x) = -5.0x_1^2 - 3.3x_2^2 + 0.2620$$

$$g_2(x) = -2.5x_1^2 - 2.0x_2^2 + 5.0x_1 - 2.840$$

$$g_1(x) - g_2(x) = -2.500x_1^2 - 1.333x_2^2 - 5.0x_1 + 3.102 = 0$$



Bayesian classification


- Bayesian classifier for normal distributions
 - Example (equiprobable classes)

$$\mu_1 = (0, 0)^T, \mu_2 = (1, 0)^T, \Sigma_1 = \begin{pmatrix} 0.10 & 0.00 \\ 0.00 & 0.15 \end{pmatrix}, \boxed{\Sigma_2 = \begin{pmatrix} 0.15 & 0.00 \\ 0.00 & 0.10 \end{pmatrix}}$$

$$g_1(x) = -5.0x_1^2 - 3.3x_2^2 + 0.2620$$

$$g_2(x) = -3.333x_1^2 - 5.0x_2^2 + 6.667x_1 - 3.071$$

$$g_1(x) - g_2(x) = -1.667x_1^2 + 1.667x_2^2 - 6.667x_1 + 3.333 = 0$$



Bayesian classification

- Bayesian classifier for normal distributions

- Classes with the same covariance matrix: decision hyperplanes

- If the classes have the same covariance matrix ($\Sigma = \Sigma_i$) then the **quadratic term** and **part of the constant term** coincide in all discrimination functions:

$$g_i(x) = -\frac{1}{2}x^T \Sigma^{-1} x + \frac{1}{2}x^T \Sigma^{-1} \mu_i + \frac{1}{2}\mu_i^T \Sigma^{-1} x - \frac{1}{2}\mu_i^T \Sigma^{-1} \mu_i + \ln p(\omega_i) + c$$

- Therefore, they disappear from equations $g_i(x) - g_j(x) = 0$.

This allows us to define **more useful and simpler discrimination functions**:

$$g_i(x) = w_i^T x + w_{i0} = \alpha_i x_1 + \beta_i x_2 + \gamma_i \quad (\text{case of 2 features})$$

$$w_i^T = \mu_i^T \Sigma^{-1}, \quad w_{i0} = \ln p(\omega_i) - \frac{1}{2}\mu_i^T \Sigma^{-1} \mu_i$$

- In this way, the discrimination functions are linear (and not quadratic) and the decision rule turns out to be a **decision hyperplane**: (2D) a straight line, (3D) a plane, ...

LINEAR CLASSIFIER

$$\rightarrow g_i(x) > g_j(x), \forall j \quad p(\omega_i|x) > p(\omega_j|x), \forall j \Rightarrow x \rightarrow \omega_i$$

$$\rightarrow g_1(x) - g_2(x) = (\alpha_1 - \alpha_2)x_1 + (\beta_1 - \beta_2)x_2 + (\gamma_1 - \gamma_2) > 0 \Rightarrow x \rightarrow \omega_1$$

- Let us have a look at **two cases of the covariance matrix**: (1) $\Sigma = \sigma^2 I$ and (2) any Σ

Bayesian classification

- Bayesian classifier for normal distributions

- Classes with the **same covariance matrix**: $\Sigma = \sigma^2 I$

- Then, **the discrimination functions** take the following form:

$$g_i(x) = \frac{1}{\sigma^2} \mu_i^T x + w_{i0} = \frac{1}{\sigma^2} \mu_i^T x + \ln p(\omega_i) - \frac{1}{2\sigma^2} \mu_i^T \mu_i$$

so that the **decision rules** can be written as:

$$g_{ij}(x) \equiv g_i(x) - g_j(x) = 0$$

$$\Rightarrow \frac{1}{\sigma^2} (\mu_i^T - \mu_j^T) x + \ln p(\omega_i) - \ln p(\omega_j) - \frac{1}{2\sigma^2} (\mu_i^T \mu_i - \mu_j^T \mu_j) = 0$$

$$\Rightarrow (\mu_i - \mu_j)^T x + \sigma^2 \ln \left(\frac{p(\omega_i)}{p(\omega_j)} \right) - \frac{1}{2} (\mu_i - \mu_j)^T (\mu_i + \mu_j) = 0$$

$$\Rightarrow (\mu_i - \mu_j)^T \left[x + \sigma^2 \ln \left(\frac{p(\omega_i)}{p(\omega_j)} \right) \frac{\mu_i - \mu_j}{\|\mu_i - \mu_j\|^2} - \frac{1}{2} (\mu_i + \mu_j) \right] = 0$$

$$\Rightarrow w^T (x - x_0) = 0$$

$$w = \mu_i - \mu_j, x_0 = \frac{1}{2} (\mu_i + \mu_j) - \sigma^2 \ln \left(\frac{p(\omega_i)}{p(\omega_j)} \right) \frac{\mu_i - \mu_j}{\|\mu_i - \mu_j\|^2}$$

Bayesian classification

- Bayesian classifier for normal distributions

- Classes with the same covariance matrix: $\Sigma = \sigma^2 I$

- Decision rule

$$g_{ij}(x) : w^T(x - x_0) = 0$$

$$w = \mu_i - \mu_j, x_0 = \frac{1}{2}(\mu_i + \mu_j) - \sigma^2 \ln \left(\frac{p(\omega_i)}{p(\omega_j)} \right) \frac{\mu_i - \mu_j}{\|\mu_i - \mu_j\|^2}$$

- Two-feature case

$$g_{ij}(x) : w^T(x - x_0) = 0$$

$$\Rightarrow (\mu_{1,1} - \mu_{2,1})x_1 + (\mu_{1,2} - \mu_{2,2})x_2 \\ - (\mu_{1,1} - \mu_{2,1})x_{0,1} - (\mu_{1,2} - \mu_{2,2})x_{0,2} = 0$$

$$\Rightarrow \alpha_{ij}x_1 + \beta_{ij}x_2 + \gamma_{ij} = 0$$

- Which is this straight line?

Bayesian classification

- Bayesian classifier for normal distributions


- Classes with the **same covariance matrix**: $\Sigma = \sigma^2 I$

- Decision rule for the **two-feature case**

$$g_{ij}(x) : w^T(x - x_0) = 0$$



$$w = \mu_i - \mu_j, x_0 = \frac{1}{2}(\mu_i + \mu_j) - \sigma^2 \ln \left(\frac{p(\omega_i)}{p(\omega_j)} \right) \frac{\mu_i - \mu_j}{\|\mu_i - \mu_j\|^2}$$

- Any point x such that $x - x_0$ is orthogonal to $\mu_i - \mu_j$ belongs to the straight line
 - x_0 is always along the vector $\mu_i - \mu_j$
 - if $p(\omega_i) = p(\omega_j)$, x_0 is the middle point between μ_i and μ_j
 - if $p(\omega_i) < p(\omega_j)$, x_0 moves towards μ_i along vector $\mu_i - \mu_j$



Bayesian classification

- Bayesian classifier for normal distributions
 - Classes with the **same covariance matrix**: $\Sigma = \sigma^2 I$
 - the circles expand to $3\sigma \equiv 98\%$



Bayesian classification

- Bayesian classifier for normal distributions
 - Classes with the **same covariance matrix**: any Σ
 - We recover the original linear discrimination functions:

$$g_i(x) = w_i^T x + w_{i0}$$

$$w_i^T = \mu_i^T \Sigma^{-1}, \quad w_{i0} = \ln p(\omega_i) - \frac{1}{2} \mu_i^T \Sigma^{-1} \mu_i$$

- Then:

$$g_{ij}(x) \equiv g_i(x) - g_j(x) = 0$$

$$\begin{aligned} & \mu_1^T \Sigma^{-1} \mu_1 - \mu_2^T \Sigma^{-1} \mu_2 + \mu_1^T \Sigma^{-1} \mu_2 - \mu_2^T \Sigma^{-1} \mu_1 \\ &= (\mu_1 - \mu_2)^T \Sigma^{-1} \mu_1 + (\mu_1 - \mu_2)^T \Sigma^{-1} \mu_2 \\ &= (\mu_1 - \mu_2)^T \Sigma^{-1} (\mu_1 + \mu_2) \end{aligned}$$

$$\Rightarrow (\mu_i - \mu_j)^T \Sigma^{-1} x + \ln \left(\frac{p(\omega_i)}{p(\omega_j)} \right) - \frac{1}{2} (\mu_i - \mu_j)^T \Sigma^{-1} (\mu_i + \mu_j) = 0$$

$$\Rightarrow (\mu_i - \mu_j)^T \Sigma^{-1} \left[x + \ln \left(\frac{p(\omega_i)}{p(\omega_j)} \right) \frac{\mu_i - \mu_j}{(\mu_i - \mu_j)^T \Sigma^{-1} (\mu_i - \mu_j)} - \frac{1}{2} (\mu_i + \mu_j) \right] = 0$$

$$\Rightarrow w^T (x - x_0) = 0$$

$$w = \Sigma^{-1} (\mu_i - \mu_j), x_0 = \frac{1}{2} (\mu_i + \mu_j) - \ln \left(\frac{p(\omega_i)}{p(\omega_j)} \right) \frac{\mu_i - \mu_j}{(\mu_i - \mu_j)^T \Sigma^{-1} (\mu_i - \mu_j)}$$


Bayesian classification

- Bayesian classifier for normal distributions
 - Classes with the same covariance matrix: any Σ

$$g_{ij} : w^T(x - x_0) = 0$$

$$w = \Sigma^{-1}(\mu_i - \mu_j)$$

$$x_0 = \frac{1}{2}(\mu_i + \mu_j) - \ln\left(\frac{p(\omega_i)}{p(\omega_j)}\right) \frac{\mu_i - \mu_j}{(\mu_i - \mu_j)^T \Sigma^{-1}(\mu_i - \mu_j)}$$



- The decision hyperplane is no longer necessarily orthogonal to $\mu_i - \mu_j$ but to $\Sigma^{-1}(\mu_i - \mu_j)$
 - $\Sigma^{-1}(\mu_i - \mu_j)$ is the result of transforming $(\mu_i - \mu_j)$ through the matrix Σ^{-1}
- x_0 is always on the vector $\mu_i - \mu_j$
 - if $p(\omega_i) = p(\omega_j)$, x_0 is the average of μ_i y μ_j
 - if $p(\omega_i) < p(\omega_j)$, x_0 moves towards μ_i along $\mu_i - \mu_j$

Bayesian classification

- Bayesian classifier for normal distributions

- Example In a two-dimensional classification problem with two equiprobable classes, the classes follow two normal distributions with the following parameters:

$$\mu_1 = (0, 0)^T, \mu_2 = (3, 3)^T, \Sigma = \Sigma_1 = \Sigma_2 = \begin{pmatrix} 1.1 & 0.3 \\ 0.3 & 1.9 \end{pmatrix}$$

Classify the sample $x = (1.0, 2.2)^T$ using a Bayesian classifier.

$$g_{12}(x) = w^T(x - x_0)$$

$$w = \Sigma^{-1}(\mu_1 - \mu_2)$$

$$x_0 = \frac{1}{2}(\mu_1 + \mu_2) - \ln\left(\frac{p(\omega_1)}{p(\omega_2)}\right) \frac{\mu_1 - \mu_2}{(\mu_2 - \mu_1)^T \Sigma^{-1} (\mu_1 - \mu_2)}$$

$$\begin{aligned} g_{12} &= (-3, -3) \begin{pmatrix} 1.1 & 0.3 \\ 0.3 & 1.9 \end{pmatrix}^{-1} \left((1, 2.2)^T - (1.5, 1.5)^T \right) \\ &= 0.36 > 0 \end{aligned}$$

Therefore, $x \rightarrow \omega_1$.

Bayesian classification

- Bayesian classifier for normal distributions

- So far we have considered **two-class cases** and derived the boundaries between class pairs through the discrimination functions:





$$g_{ij}(x) = g_i(x) - g_j(x) = 0$$

- For **two classes** if $g_{12}(x) > 0$, then $x \rightarrow \omega_1$ [if $g_1(x) > g_2(x)$, then $x \rightarrow \omega_1$]
if $g_{12}(x) < 0$, then $x \rightarrow \omega_2$ [if $g_1(x) < g_2(x)$, then $x \rightarrow \omega_2$]

- If there are more than 2 classes, we have to determine $g_i \mid g_i > g_j, \forall j \neq i$.

In case of using the joint discrimination functions g_{ij} to decide, we have to consider several g_{ij} . For example, for 3 classes:

$$\begin{array}{c} \mu_i \\ \hline \hline (0,0) \\ (-2,2) \\ (+2,2) \\ \hline \hline \sigma = 1 \end{array}$$



- to determine R_i we need several g_{ik}, g_{ki}
M classes: M-1 pairs

Bayesian classification

- Bayesian classifier for normal distributions

- Minimum distance classifiers

- We can see the above from another point of view
 - We assume **equiprobable classes with the same covariance matrix**. Then:

$$\left. \begin{aligned} g_i(x) &= c_i - \frac{1}{2}(x - \mu_i)^T \Sigma_i^{-1} (x - \mu_i) + \ln p(\omega_i) \\ c_i &= -\frac{L}{2} \ln 2\pi - \frac{1}{2} \ln |\Sigma_i| \end{aligned} \right\} \rightarrow$$

$$\rightarrow g_i(x) = -\frac{1}{2}(x - \mu_i)^T \Sigma^{-1} (x - \mu_i)$$

- We assign \mathbf{x} to the class for which the probability is greater $\Rightarrow g_i(x) > g_j(x) \forall j \neq i$
 - (1) If $\Sigma = \sigma^2 \mathbf{I}$, $g_i(x)$ is higher the closer it is x to μ_i
 \Rightarrow assign \mathbf{x} to the class whose center μ_i is closer (**Euclidean distance**)
 - (2) For generic Σ , we have to assign \mathbf{x} to the class for which the following expression takes a lowest value:
 $d_m^2 = (x - \mu_i)^T \Sigma^{-1} (x - \mu_i)$
 - $d_m \equiv$ **Mahalanobis distance** (considers the scattering present in the features)

Bayesian classification

- Bayesian classifier for normal distributions

- Example In a two-dimensional classification problem into two equiprobable classes the classes have two normal distributions with the following parameters:

$$\mu_1 = (0, 0)^T, \mu_2 = (3, 3)^T, \Sigma = \Sigma_1 = \Sigma_2 = \begin{pmatrix} 1.1 & 0.3 \\ 0.3 & 1.9 \end{pmatrix}$$

Classify the vector $\mathbf{x} = (1.0, 2.2)^T$ using a Bayesian classifier.

$$\begin{aligned} d_m^2(\mathbf{x}, \mu_1) &= (\mathbf{x} - \mu_1)^T \Sigma^{-1} (\mathbf{x} - \mu_1) \\ &= (1.0, 2.2) \begin{pmatrix} 0.95 & -0.15 \\ -0.15 & 0.55 \end{pmatrix} \begin{pmatrix} 1.0 \\ 2.2 \end{pmatrix} = 2.952 \end{aligned}$$

$$\begin{aligned} d_m^2(\mathbf{x}, \mu_2) &= (\mathbf{x} - \mu_2)^T \Sigma^{-1} (\mathbf{x} - \mu_2) \\ &= (-2.0, -0.8) \begin{pmatrix} 0.95 & -0.15 \\ -0.15 & 0.55 \end{pmatrix} \begin{pmatrix} -2.0 \\ -0.8 \end{pmatrix} = 3.672 \end{aligned}$$

- $d_m(\mathbf{x}, \mu_1) < d_m(\mathbf{x}, \mu_2) \Rightarrow \mathbf{x} \rightarrow \omega_1$.
- REMARK: the Euclidean distances would be $d_e(\mathbf{x}, \mu_1) = 2.417$ and $d_e(\mathbf{x}, \mu_2) = 2.154$, so, if we used them, we would assign $\mathbf{x} \rightarrow \omega_2$.

Contents

- Introduction
- Bayesian classification
- Estimation of probability density functions
- Linear discriminant functions and the perceptron algorithm

Estimation of probability density functions

- **Estimation of probability density functions (pdfs)**
 - The Bayesian classifier assumes that we have **knowledge on the pdfs of the classes** of the problem
 - There are different methods to get this type of information:
 - The expression of the pdf is known but the parameters are unknown
→ **parametric estimation**
 - **Maximum likelihood estimators**
 - others
 - The expression of the pdf is not known → **non-parametric estimation**
 - **Parzen windows** method
 - **k-nearest neighbours** method (KNN)
 - others

Estimation of probability density functions

- **Estimation of probability density functions**
 - **Maximum likelihood estimators**
 - Let us consider an M -class classification problem whose samples are distributed in accordance to $p(x|\omega_i; \theta_i)$, $i = 1, \dots, M$, where θ_i is the **vector of parameters** for class ω_i
 - It's about estimating θ_i by means of a set of samples x_1, x_2, \dots, x_k from class ω_i
 - We assume that the samples of one class do not affect the estimation of parameters for the other classes in order to formulate the problem irrespective of the class
 - ⇒ the estimation is repeated for each class
 - In this way, given the statistically independent samples x_1, x_2, \dots, x_N from $p(x_k|\theta)$, we calculate the following joint pdf:

$$p(X; \theta) = p(x_1, x_2, \dots, x_N; \theta) = \prod_{k=1}^N p(x_k; \theta)$$

- Then, the **maximum likelihood estimator** of θ is given by:

$$\theta_{ML} \mid p(X; \hat{\theta}_{ML}) = \max\{p(X; \theta)\}$$

- This represents that θ that better explains samples x_1, x_2, \dots, x_N

Estimation of probability density functions

- Estimation of probability density functions
 - Maximum likelihood estimators
 - To simplify the calculations we will go once again to the function $\ln(\cdot)$ to define the **log-likelihood function**:

$$L(\theta) \equiv \ln \prod_{k=1}^N p(x_k; \theta) = \sum_{k=1}^N \ln p(x_k; \theta)$$

- Now, you can find the derivative of the log-likelihood and equal to 0:

$$\frac{\partial L(\theta)}{\partial \theta} = \sum_{k=1}^N \frac{\partial \ln p(x_k; \theta)}{\partial \theta} = \sum_{k=1}^N \frac{1}{p(x_k; \theta)} \frac{\partial p(x_k; \theta)}{\partial \theta} = 0$$

- For sufficiently high N values, the maximum likelihood estimator is **asymptotically unbiased**, follows a **normal distribution** and exhibits **minimal variance**

Estimation of probability density functions

- Estimation of probability density functions
 - Maximum likelihood estimators
 - L-dimensional Gaussian distribution with Σ known

$$\hat{\mu}_{ML} = \frac{1}{N} \sum_{k=1}^N x_k$$


- L-dimensional Gaussian distribution, μ and Σ unknown

$$\hat{\mu}_{ML} = \frac{1}{N} \sum_{k=1}^N x_k, \quad \hat{\Sigma}_{ML} = \frac{1}{N} \sum_{k=1}^N (x_k - \hat{\mu}_{ML})(x_k - \hat{\mu}_{ML})^T$$

Estimation of probability density functions

- **Estimation of probability density functions**
 - Non-parametric estimation: **first approximation**
 - It is about estimating a certain pdf $p(x)$ without setting any expression for the pdf
 - Let us assume we have N independent samples x_1, x_2, \dots, x_N that come from the pdf that we want to estimate
 - To this end, we build a **histogram** using bins R_j of the same size:

for different sizes of R_j




- $k_{N,Rj}$ = how many of the N samples belong to bin R_j
- $\{k_{N,Rj}/N\}$ is an approximation of P_{Rj} = probability that x belongs to R_j :

$$P_{R_j} = p(x \in R_j) = \int_{R_j} p(x') dx'$$

Estimation of probability density functions

- **Estimation of probability density functions**
 - Non-parametric estimation: **first approximation**
 - We now define sufficiently small regions R around x , aiming at calculating $p(x)$. In this regard, if $p(x)$ is assumed constant inside R :
$$P_R = \int_R p(x') dx' \approx p(x) \int_R dx' = p(x)V$$
 - V is the (hyper)volume occupied by region R (1D – length, 2D – area, 3D – volume, etc.)
 - e.g. if R is a (hyper)cube of dimension L and side length h , $V = h^L$
 - Therefore: $p(x) \approx \frac{P_R}{V} = \frac{k_{N,R}/N}{V} = p_{N,R}(x)$
 - $p_{N,R}(x) \rightarrow p(x)$ as $N \rightarrow \infty$ if the following holds:
 - $V \rightarrow 0$ (small regions)
 - $k_{N,R} \rightarrow \infty$ (sufficient number of samples in each R)
 - $k_{N,R}/N \rightarrow 0$ (high total number of samples)
 - **In short:** for each x , $p(x)$ is approximated by defining a small R region around x and counting how many x_i fall into R ($= k_{N,R}$); if $k_{N,R} \uparrow\uparrow$ and $N \uparrow\uparrow\uparrow$, then $p_{N,R}(x) \approx p(x)$



Estimation of probability density functions


- **Estimación de funciones de densidad de probabilidad**
 - Non-parametric estimation: **Parzen windows** (Parzen, 1962)
 - Let us consider a region R shaped like a (hyper) L -dimensional cube of side h_N . Then:
$$V_N = (h_N)^L$$
 - Let us consider the following function (**box function** or **kernel**):
$$\varphi(u) = \begin{cases} 1 & |u_j| \leq \frac{1}{2}, j = 1, \dots, L \\ 0 & \text{otherwise} \end{cases}$$
 - So, for a certain x :
$$\varphi\left(\frac{x-x_i}{h_N}\right) = 1 \text{ if } x_i \in \text{(hyper)cube with volume } V_N \text{ centered at } x$$
 - Then, the number of samples that are inside the x -centered (hyper)cube is:

$$k_N(x) = \sum_{i=1}^N \varphi\left(\frac{x - x_i}{h_N}\right)$$

Estimation of probability density functions

- Estimation of probability density functions
 - Non-parametric estimation: **Parzen windows**

- At last:

$$p_N(x) = \frac{k_N(x)/N}{V_N} = \frac{1}{N} \sum_{i=1}^N \frac{1}{h_N^L} \varphi\left(\frac{x - x_i}{h_N}\right)$$


- Therefore, given a certain x , to obtain the estimate of $p(x)$: (1D case)

```
def parzen_box_1D(x, X, h):
    # x = point where to evaluate the PDF
    # X = table of samples
    # h = side of the (hyper)cube

    N = X.shape[0]; kn = 0
    for i in range(N):
        if abs((x - X[i])/h) <= 0.5:
            kn = kn+1
    p = (kn/N)/h
    return p
```

Estimation of probability density functions

- **Estimation of probability density functions**
 - Non-parametric estimation: **Parzen windows**
 - Comments on p_N :
 - p_N is a legitimate pdf

1. $p_N(x) \geq 0, \forall x$ (obvious, it is a sum of 1's)
2. $\int p_N(x)dx = 1$

$$\begin{aligned} & \int \frac{1}{N} \sum_{i=1}^N \frac{1}{h_N^L} \varphi \left(\frac{x - x_i}{h_N} \right) dx = \frac{1}{N} \sum_{i=1}^N \frac{1}{h_N^L} \int \varphi \left(\frac{x - x_i}{h_N} \right) dx = \\ & = \{u = x - x_i\} = \frac{1}{N} \sum_{i=1}^N \frac{1}{h_N^L} \int \varphi \left(\frac{u}{h_N} \right) du = \\ & = \frac{1}{N} \sum_{i=1}^N \frac{1}{h_N^L} \int_{(-h_N/2, \dots, -h_N/2)}^{(+h_N/2, \dots, +h_N/2)} du = \frac{1}{N} \sum_{i=1}^N \frac{1}{h_N^L} h_N^L = 1 \end{aligned}$$

Estimation of probability density functions

- Estimation of probability density functions
 - Non-parametric estimation: **Parzen windows**


- Comments on p_N :

- function $p_N(x)$ can be seen as the average of N functions centred in the samples x_i

$$p_N(x) = \frac{1}{N} \sum_{i=1}^N \frac{1}{h_N^L} \varphi\left(\frac{x - x_i}{h_N}\right)$$


$$\left\{ \begin{array}{l} p_N(x) = \frac{1}{N} \sum_{i=1}^N \delta_N(x - x_i) \\ \delta_N(u) = \frac{1}{h_N^L} \varphi\left(\frac{u}{h_N}\right) \end{array} \right.$$

- if h_N is **large**, the amplitude of δ_N is **small** and p_N is the superposition of N *wide pulses*
- if h_N is **small**, the amplitude of δ_N is **large** and p_N is the superposition of N *narrow pulses*
- $h_N \rightarrow 0 \Rightarrow \delta_N \rightarrow \delta$



Estimation of probability density functions

- Estimation of probability density functions
 - Non-parametric estimation: **Parzen windows**
 - Example: N random values extracted from a distribution $N(0,1)$ – rectangular kernel



Estimation of probability density functions

- **Estimation of probability density functions**

- Non-parametric estimation: **Parzen windows**

- As we have already seen in the previous example, when approximating continuous functions $[p(\cdot)]$ by discontinuous kernel functions $[\varphi(\cdot)]$, the resulting estimate also presents **discontinuities**
 - To avoid this, it is suggested to use **continuous kernels** $\varphi(\cdot)$
 - It can be shown that the resulting estimate $p_N(x)$ is a legitimate pdf if:

$$\varphi(u) \geq 0 \text{ and } \int \varphi(u)du = 1$$


- One of the most commonly used kernels is the Gaussian kernel (mean 0, variance 1): $\varphi(u) = \frac{1}{\sqrt{(2\pi)^L}} e^{-\frac{1}{2}u^T u}$

$$\begin{aligned} p_N(x) &= \frac{1}{N} \sum_{i=1}^N \frac{1}{h_N^L} \varphi \left(\frac{x - x_i}{h_N} \right) \\ &= \frac{1}{N h_N^L} \sum_{i=1}^N \frac{1}{\sqrt{(2\pi)^L}} e^{-\frac{1}{2}u_i^T u_i} \\ &\quad \left(u_i = \frac{x - x_i}{h_N} \right) \end{aligned}$$


```
def parzen_gauss_1D(x, X, h):
    # x = point where to evaluate the PDF
    # X = data samples
    # h = side of the (hyper)cube
    N = X.shape[0]; kn = 0
    for i in range(N):
        kn += 1/(sqrt(2*pi))*exp(-0.5*((x-X[i])/h)**2)
    p = (kn/N)/h
    return p
```

Estimation of probability density functions

- Estimation of probability density functions
 - Non-parametric estimation: **Parzen windows**
 - Now $p_N(x)$ is obtained as the average of N Gaussians centered on the samples x_i




- Effect of varying bandwidth h_N



Estimation of probability density functions

- Estimation of probability density functions
 - Non-parametric estimation: **Parzen windows**
 - Example: N random values extracted from a distribution $N(0,1)$ – Gaussian kernel



Estimation of probability density functions

- **Estimation of probability density functions**
 - Non-parametric estimation: **k nearest neighbours**
 - Given x and a collection of samples x_1, x_2, \dots, x_N from a certain pdf $p(x)$, to estimate $p(x)$:
 - **Parzen windows method** – first, a search volume is set around x , V_N , and next we determine the number of samples k_N belonging to that volume

We set V_N and find $k_N \rightarrow p_N(x) = \frac{k_N/N}{V_N}$


- **Method of the k nearest neighbors** – we first find the k_N samples nearest to x , and next we determine the minimal volume V_{kN} where they are contained in

We set k_N and find $V_N \rightarrow p_N(x) = \frac{k_N/N}{V_{kN}}$

```
def knn_1D(x, X, k):
    # x = point where to evaluate the PDF
    # X = data samples
    # k = number of neighbours
    N = X.shape[0]; d = []
    for i in range(N):
        d.append(abs(x-X[i]))
    d.sort()
    v = 2 * d[min(N, k)-1]
    p = (k/N)/v
    return p
```

Estimation of probability density functions

- Estimation of probability density functions
 - Non-parametric estimation: **k nearest neighbours**
 - Example: N random values sampled from a distribution $N(0,1)$



- In this case, it has been used: $k_N = \sqrt{N}$
- In the case of Parzen windows, it is suggested to use: $h_N = \frac{h_1}{\sqrt{N}}$

Contents

- Introduction
- Bayesian classification
- Estimation of probability density functions
- Linear discriminant functions and the perceptron algorithm

Linear discrimination functions and the perceptron algorithm

- We have already seen that, depending on the pdf's of the classes (Gaussian case), a Bayesian classifier can derive in a set of **linear discrimination functions**. For example, for 2 classes:


$$g_{12}(x) = w^T(x - x_0), \quad g_{12}(x) = \begin{cases} > 0 & x \in \omega_1 \\ < 0 & x \in \omega_2 \end{cases}$$

- simple and computationally very interesting classifier
- In this section, we concentrate again on linear discrimination functions, but from a different perspective: **we do not assume any pdf for the classes**
 - Therefore, regardless of the pdf of the classes, we expect them to be separable by (hyper)planes (1D – point, 2D – straight line, 3D – plane, etc.)
 - In this case it is said that the **classes are linearly separable**
 - We will see how you can find a (hyper)plane that separates the classes from each other (**perceptron algorithm**)

Linear discrimination functions and the perceptron algorithm

- Linear discrimination functions


- Goal: find a (hyper)plane that allows us to separate the training samples in 2 classes



$$g_{12}(x) = \mathbf{w}^T(\mathbf{x} - \mathbf{x}_0), \quad g_{12}(x) = \begin{cases} > 0 & x \in \omega_1 \\ < 0 & x \in \omega_2 \end{cases}$$

(hyper)plane: $\mathbf{x} \mid g_{12}(\mathbf{x}) = \mathbf{w}^T(\mathbf{x} - \mathbf{x}_0) = 0$

- For example, for $L = 2$ features:



$$(\mathbf{w}_1, \mathbf{w}_2) \left[\begin{pmatrix} x_1 \\ x_2 \end{pmatrix} - \begin{pmatrix} x_{01} \\ x_{02} \end{pmatrix} \right] = 0$$



$$w_1 x_1 + w_2 x_2 - (w_1 x_{01} + w_2 x_{02}) = 0$$

$$w_1 x_1 + w_2 x_2 + w_0 = 0 \quad (\text{normal form})$$

$$\begin{aligned} x_2 &= \left(-\frac{w_1}{w_2} \right) x_1 + \left(-\frac{w_0}{w_2} \right) && (\text{slope-intercept form}) \\ &= ax_1 + b \end{aligned}$$

Linear discrimination functions and the perceptron algorithm

- **Linear discrimination functions.** For now, we will only consider hyperplanes which go through the origin



- This means that $x_0 = (0,0)^T$ and thus in $w^T(x - x_0)$
 $w_0 = w^T x_0 = 0$

\Downarrow

$$w_1 x_1 + w_2 x_2 = 0$$

\Downarrow

$$w^T x = 0$$

- For any point x_i outside the hyperplane we can state:

$$w^T x_i = \|w\| \|x_i\| \cos \beta \neq 0$$

\Downarrow $\|w\| = 1$

$$\|x_i\| \cos \beta = d$$


- The sign of d depends on the **relative position** of x_i with regard to the hyperplane:

$$\begin{cases} \beta \in [-\frac{\pi}{2}, +\frac{\pi}{2}] & w^T x_i = d > 0 \\ |\beta| > \frac{\pi}{2} & w^T x_i = d < 0 \end{cases}$$

- $|d| = |w^T x_i|$ indicates **how far away the sample is** from the (hyper)plane of discrimination

Linear discrimination functions and the perceptron algorithm

- We assume that the classes ω_1 and ω_2 are **linearly separable**, i.e. the hyperplane exists
- The goal is thus to find a function $g_{12}(x) = w^T x$ such that $g_{12}(x)$ is as follows for each sample x_i of the training set:



$$g_{12}(x_i) = w^T x_i > 0, \forall x_i \in \omega_1$$

$$g_{12}(x_i) = w^T x_i < 0, \forall x_i \in \omega_2$$

– g_{12} defined in this way is also named as a **discrimination function**, which in this case turns out to be linear, and thus it is a **linear discrimination function**

- To find the hyperplane, we consider the following function (**perceptron cost**):

$$J(w) = \sum_{x_i \in \mathcal{Y}} (\delta_{x_i} w^T x_i)$$

where: – \mathcal{Y} is the set of samples x_i wrongly classified by w

- $\delta_{x_i} = -1$ if $x_i \in \omega_1$ y $\delta_{x_i} = +1$ if $x_i \in \omega_2$
- $J(w) \geq 0, \forall w$ ($x_i \in \omega_1$ but $x_i \rightarrow \omega_2$, then $\delta_{x_i} \omega^T x_i = (-1)(< 0) > 0$)
- $J(w) = 0$ if all samples are well classified ($\mathcal{Y} = \emptyset$)
- $J(w)$ is piece-wise linear \Rightarrow minimization is not trivial

Linear discrimination functions and the perceptron algorithm

- The **perceptron algorithm** (Rosenblatt, 1950s) is able to find, through the next iterative approach, the required hyperplane: (sort of **gradient descent**)

$$w(t+1) = w(t) - \rho_t \sum_{x_i \in \mathcal{Y}} \delta_{x_i} x_i$$

$$w(0) = \text{any vector of } \mathbb{R}^L$$

- The algorithm converges if the classes are **linearly separable** and if the sequence of learning factor values ρ_t meets certain conditions:

$$\lim_{t \rightarrow \infty} \sum_{k=0}^t \rho_k = \infty, \quad \lim_{t \rightarrow \infty} \sum_{k=0}^t \rho_k^2 < \infty$$

where t denotes iteration number.

- The sequence ρ_t determines the convergence speed.
- For instance, $\rho_t = c/t$ and $\rho_t = \rho$ (ρ bounded) meet those conditions.

Linear discrimination functions and the perceptron algorithm

- **Details** on the use of the (basic) perceptron algorithm:

(1a) To deal with hyperplanes that do not contain the origin, the feature vectors must be **augmented** in one additional dimension $x_i^* = (x_i, 1)^T$. In this way:

$$w^T x + w_0 = 0 \equiv \underbrace{(w_1, w_2, \dots, w_L, w_0)}_{w^T} \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_L \\ 1 \end{pmatrix} = (w^*)^T x^* = 0$$

❖ We will use w^T instead of $(w^*)^T$ and x instead of x^* to simplify the notation

(1b) The rule for modifying w within every iteration can be generically implemented as follows:

↑
1 iteration
↓

$$w(t+1) = w(t) - \rho_t \sum_{x_i \in \mathcal{Y}} \delta_{x_i} x_i \Rightarrow$$

$S = \underbrace{(0, 0, \dots, 0)}_{L+1}^T$
for $i = 1$ **to** n_training_samples
 if $x_i \in \omega_1$ **and** $w(t)^T x_i < 0$ **then** $S = S + x_i$
 if $x_i \in \omega_2$ **and** $w(t)^T x_i > 0$ **then** $S = S - x_i$
end for
 $w(t+1) = w(t) + \rho_t S$

Linear discrimination functions and the perceptron algorithm

- Details of the (basic) perceptron algorithm:

(1) Example 1

- The dashed line corresponds to:

$$(1, 1, -0.5)x =$$


$$x_1 + x_2 - 0.5 = 0$$

where $w(t) = (1, 1, -0.5)^T$ is the result of the previous step of the perceptron algorithm using $\rho_t = \rho = 0.7$

- The incorrectly classified samples are:
 $(0.40, 0.05)^T$ and $(-0.20, 0.75)^T$
- The new iteration yields

$$w(t+1) = w(t) - \rho_t \sum_{x_i \in \mathcal{Y}} \delta_{x_i} x_i$$

$$w(t+1) = \begin{pmatrix} 1 \\ 1 \\ -0.5 \end{pmatrix} - 0.7(-1) \begin{pmatrix} 0.4 \\ 0.05 \\ 1 \end{pmatrix} - 0.7(+1) \begin{pmatrix} -0.2 \\ 0.75 \\ 1 \end{pmatrix} = \begin{pmatrix} 1.42 \\ 0.51 \\ -0.5 \end{pmatrix}$$



Linear discrimination functions and the perceptron algorithm

- Details of the (basic) perceptron algorithm:

(1) Example 1

- The dashed line corresponds to:

$$(1, 1, -0.5)x =$$

$$x_1 + x_2 - 0.5 = 0$$

where $w(t) = (1, 1, -0.5)^T$ is the result of the previous step of the perceptron algorithm using $\rho_t = \rho = 0.7$


- The incorrectly classified samples are:
 $(0.40, 0.05)^T$ and $(-0.20, 0.75)^T$
- The new iteration yields

$$w(t+1) = w(t) - \rho_t \sum_{x_i \in \mathcal{Y}} \delta_{x_i} x_i$$

$$w(t+1) = \begin{pmatrix} 1 \\ 1 \\ -0.5 \end{pmatrix} - 0.7(-1) \begin{pmatrix} 0.4 \\ 0.05 \\ 1 \end{pmatrix} - 0.7(+1) \begin{pmatrix} -0.2 \\ 0.75 \\ 1 \end{pmatrix} = \begin{pmatrix} 1.42 \\ 0.51 \\ -0.5 \end{pmatrix}$$

- The resulting hyperplane classifies correctly all the samples and the algorithm ends with:

$$1.42x_1 + 0.51x_2 - 0.5 = 0$$








Linear discrimination functions and the perceptron algorithm

- Details of the (basic) perceptron algorithm:

(1) Example 2

$$\begin{aligned}\rho &= 1.5 \\ w(0) &= +0.74 & +0.67 & +0.92 \\ w(1) &= -218.29 & +77.27 & +23.42 \\ w(2) &= -152.98 & +179.38 & +11.42 \\ w(3) &= -139.43 & +188.33 & +9.92 \\ w(4) &= -125.88 & +197.29 & +8.42\end{aligned}$$



Linear discrimination functions and the perceptron algorithm

- Details of the (basic) perceptron algorithm:

(1) Implementation

```
def perceptron_2D(X, y, rho, nit):  
  
    N = X.shape[0]  
    w = np.zeros(3)  
    for t in range(nit):  
        S = np.zeros(3); ic = 0  
        for i in range(N):  
            xs = [X[i,0], X[i,1], 1]  
            if np.dot(w,xs) < 0 and y[i] == 1:      # 1 ≡ w1  
                S = S + xs; ic += 1  
            elif np.dot(w,xs) >= 0 and y[i] == 0: # 0 ≡ w2  
                S = S - xs; ic += 1  
            if ic == 0: # Y = empty set  
                break  
        else:  
            w = w + rho * S  
  
    w = w / sqrt(w[0]**2 + w[1]**2)  
    return w
```

$$w(t+1) = w(t) - \rho_t \sum_{x_i \in \mathcal{Y}} \delta_{x_i} x_i$$

Linear discrimination functions and the perceptron algorithm

- **Details** of the (basic) perceptron algorithm:

(2) The **pocket algorithm**: dealing with non-linearly separable datasets

- Stops after a number of iterations (T), providing the best hyperplane which has been found along that number of iterations
- Partially solves the **convergence problem** of the original perceptron algorithm when the classes are not linearly separable

(1) Initialize $w(0)$ randomly

(2) $w_s = w(0)$

h_s = no. samples correctly classified by $w(0)$

(3) **for** $t = 0$ **to** T

(3.1) $w(t + 1) = w(t) - \rho_t \sum_{x_i \in \mathcal{Y}} \delta_{x_i} x_i$

(3.2) h = no. samples correctly classified by $w(t + 1)$

(3.3) **if** $h > h_s$

then $w_s = w(t + 1), h_s = h$


end if

(4) **end for**

Linear discrimination functions and the perceptron algorithm

- **Classification device:**

- Once the perceptron algorithm has converged with e.g. $w = (w_1, w_2, \dots, w_L, w_0)$, then the following structure can be considered to implement the classification operation:



- The perceptron can be considered as the basic building element for more complex learning machines, e.g. **neural networks**

Lecture 3.1

Supervised learning: Bayesian and linear classifiers



Universitat
de les Illes Balears

Departament
de Ciències Matemàtiques
i Informàtica

11752 Aprendizaje Automático
11752 Machine Learning
Máster Universitario
en Sistemas Inteligentes

Alberto ORTIZ RODRÍGUEZ